

# **Establish Ingestion and Extraction patterns using Azure and Snowflake. (For Batch-time and Real-time)**

**Guided by- Amit Malik**

## **Team Members-**



**Ashwin Nair**

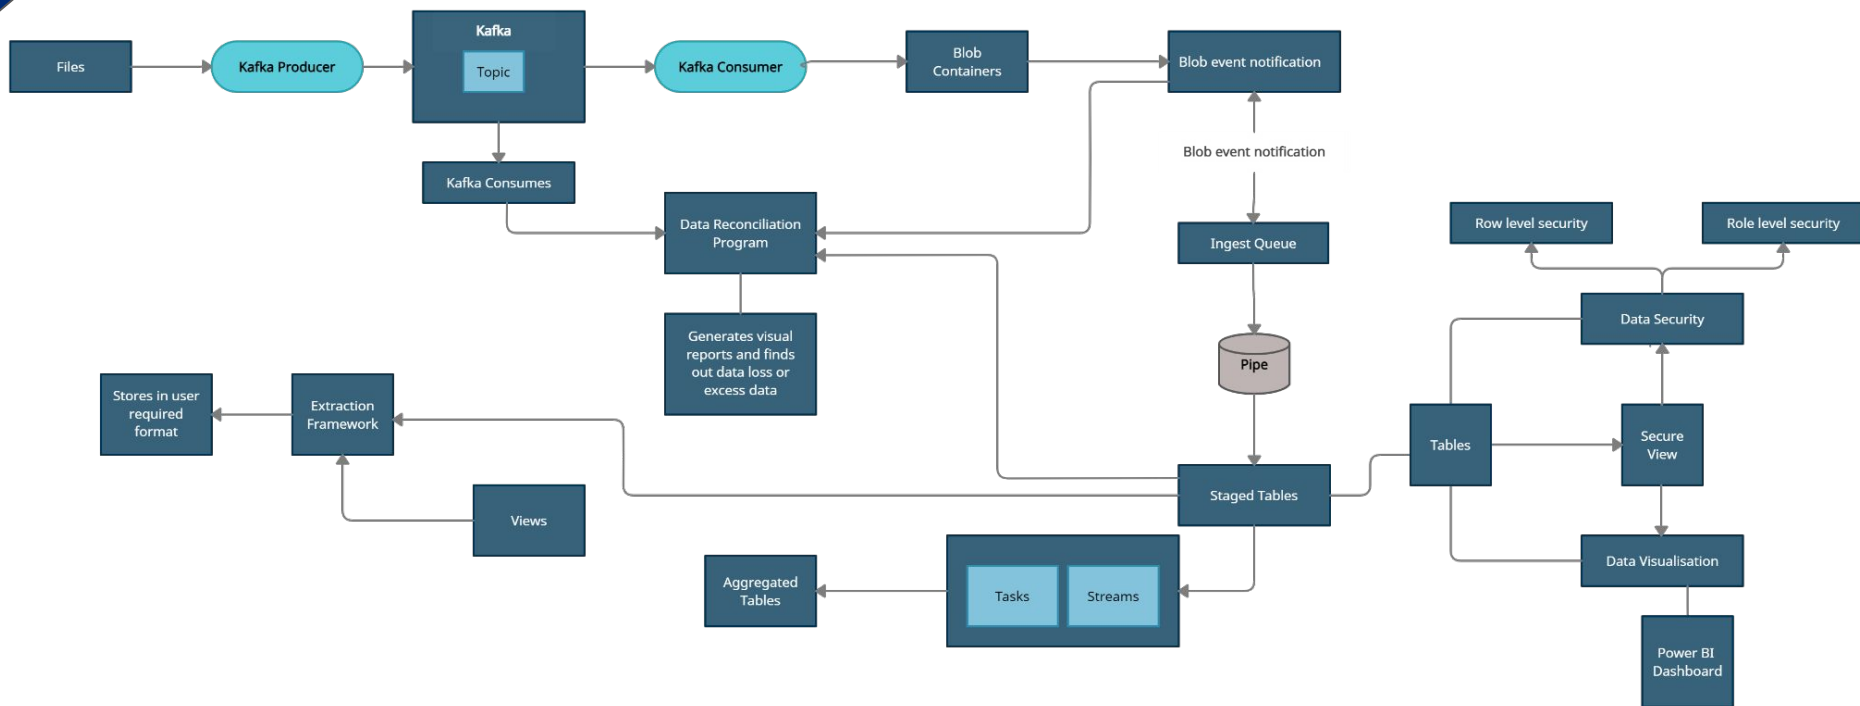


**Rajat Mishra**



**Nehal Tiwari**

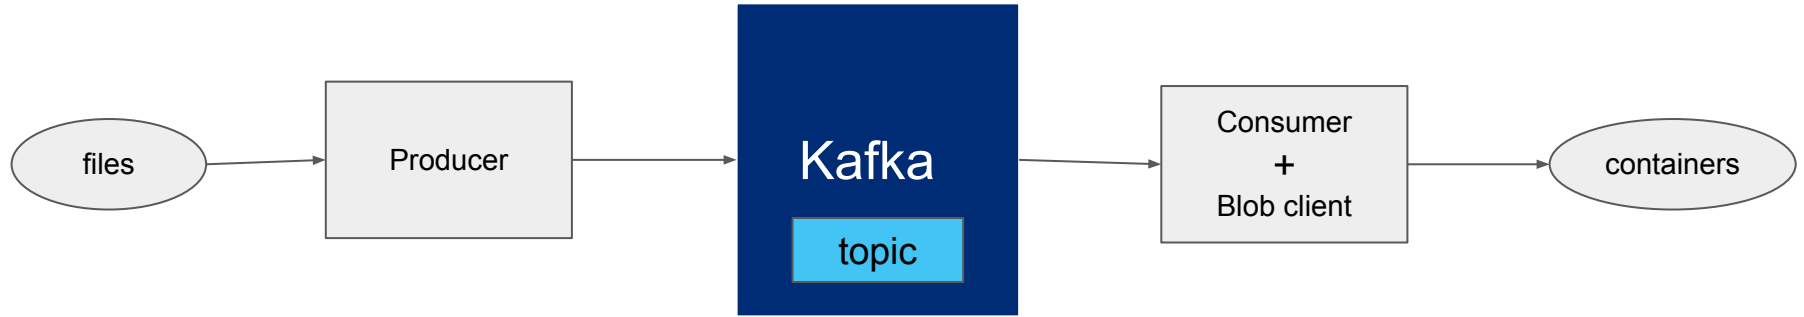
## Work-flow:



# Data Ingestion

kafka to azure

Azure blob  
containers



- Each kafka topic is mapped to a single azure storage which is in turn mapped to a snowflake database.

# Data Ingestion

## Kafka config specifies topic and maps each file to table

kafka - Notepad

File Edit Format View Help

[kafka]

topic = kafkaazure

bootstrap-servers = localhost:9092

encode = utf-8

timezone = US/Pacific

files = country\_wise\_latest.csv,dcovid\_19\_clean\_complete.csv,dday\_wise.

country\_wise\_latest.csv = COUNTRY\_WISE\_LATEST

dcovid\_19\_clean\_complete.csv = COVID\_19\_CLEAN\_COMPLETE

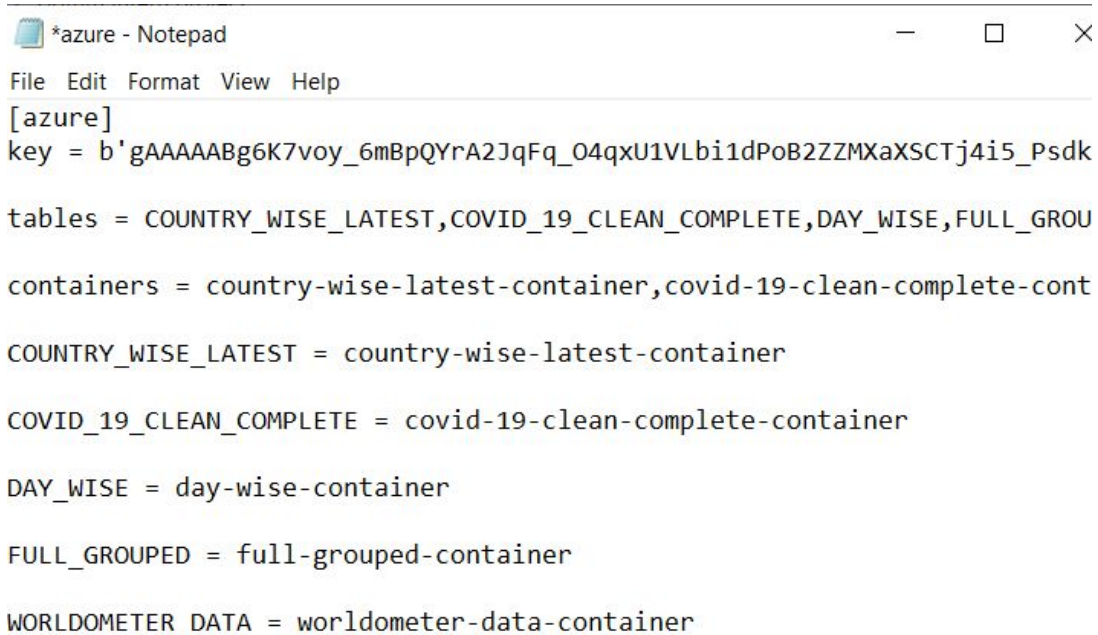
dday\_wise.csv = DAY\_WISE

dfull\_grouped.csv = FULL\_GROUPED

dworldometer\_data.csv = WORLDOMETER\_DATA

# Data Ingestion

## Setting up azure config file for consumer code to connect to Azure



```
*azure - Notepad
File Edit Format View Help
[azure]
key = b'gAAAAABg6K7voy_6mBpQYrA2JqFq_04qxU1VLbi1dPoB2ZZMXaXSCTj4i5_Psdk

tables = COUNTRY_WISE_LATEST,COVID_19_CLEAN_COMPLETE,DAY_WISE,FULL_GROU
containers = country-wise-latest-container,covid-19-clean-complete-cont

COUNTRY_WISE_LATEST = country-wise-latest-container

COVID_19_CLEAN_COMPLETE = covid-19-clean-complete-container

DAY_WISE = day-wise-container

FULL_GROUPED = full-grouped-container

WORLDOMETER_DATA = worldometer-data-container
```

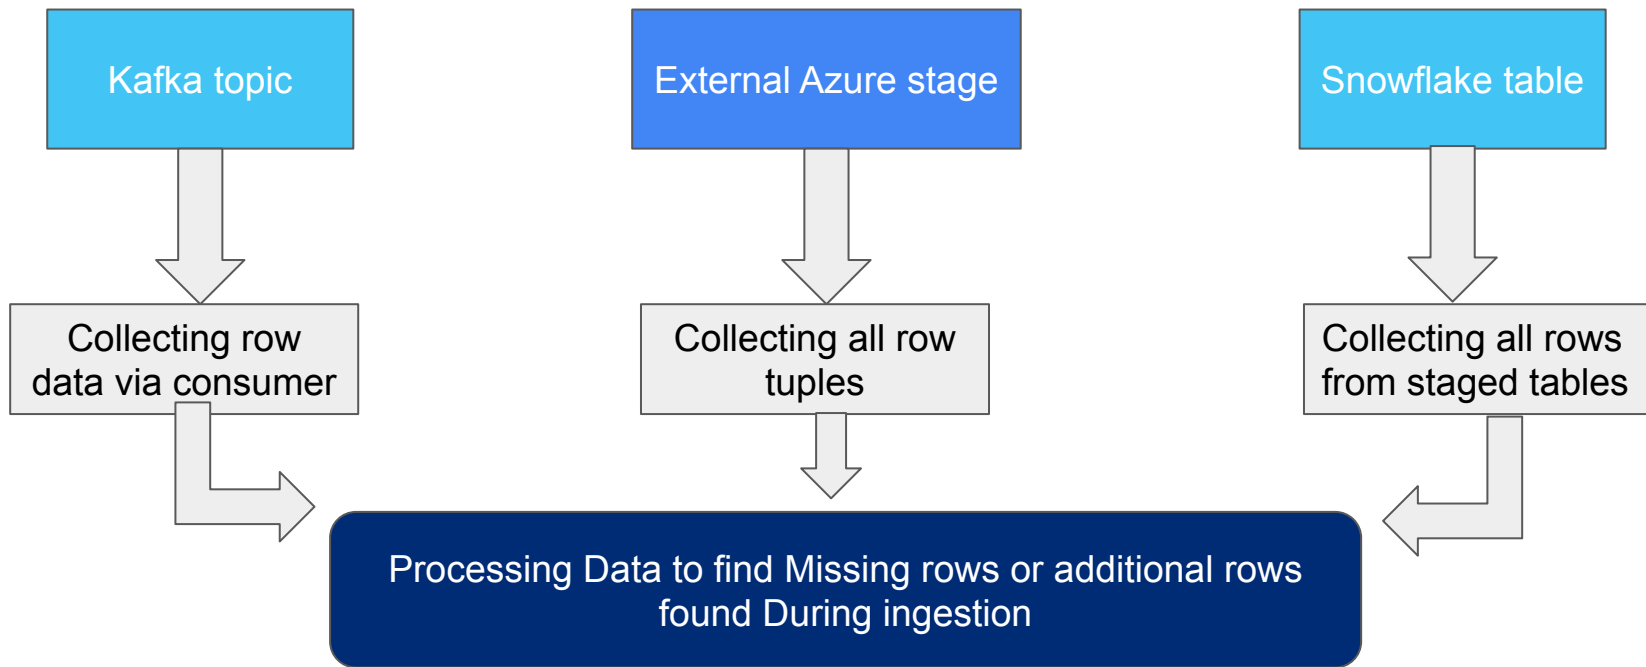
# Data Ingestion

Image depicting the mapped containers in azure storage

The screenshot shows the Microsoft Azure portal interface. The browser address bar displays the URL: <https://portal.azure.com/#@amitmaliksouthcarolinaoutlook.onmicrosoft.com/resource/subscriptions/C2ac6340-835c-40bb-8e5b-bebcead3d3/reso...>. The page title is 'ashwinstorage | Containers'. The left sidebar contains navigation options: Overview, Activity log, Tags, Diagnose and solve problems, Access Control (IAM), Data migration, Events, Storage Explorer (preview), Data storage, Containers (selected), File shares, Queues, Tables, Security + networking, Networking, Azure CDN, Access keys, and Shared access signature. The main content area shows a table of containers with columns: Name, Last modified, Public access level, and Lease state. The table lists six containers, all with 'Private' access and 'Available' lease state.

Name	Last modified	Public access level	Lease state
<input type="checkbox"/> Slogs	7/5/2021, 12:52:47 AM	Private	Available
<input type="checkbox"/> country-wise-latest-container	7/12/2021, 5:04:08 AM	Private	Available
<input type="checkbox"/> covid-19-clean-complete-container	7/12/2021, 3:56:29 AM	Private	Available
<input type="checkbox"/> day-wise-container	7/12/2021, 3:57:13 AM	Private	Available
<input type="checkbox"/> full-grouped-container	7/12/2021, 3:58:03 AM	Private	Available
<input type="checkbox"/> worldometer-data-container	7/12/2021, 3:59:24 AM	Private	Available


# Data Reconciliation




# Data Reconciliation

Accessible visuals and dataframes in python program


```
data reconciliation for DAY_WISE  
0 940 1128
```



```
1316 752 0
```



```
188 752 0
```





# Data Reconciliation

The program provides a dictionary of dictionary of tables with missing Data

Country_wise_latest	table_name	table_name	table_name	table_name
---------------------	------------	------------	------------	------------



Missing rows kafka to azure	Additional rows found in azure	Missing rows azure to snowflake	Additional rows in snowflake from other sources	Missing rows kafka to snowflake	Additional rows snowflake
-----------------------------	--------------------------------	---------------------------------	---	---------------------------------	---------------------------



Data  
representatio  
n in Tabular  
format

# Data Reconciliation

## Identification of Missing Records

```
dataset['COUNTRY_WISE_LATEST']['azure']
```

	\$1::VARCHAR	\$2::VARCHAR	\$3::VARCHAR	\$4::VARCHAR	\$5::VARCHAR	\$6::VARCHAR	\$7::VARCHAR	\$8::VARCHAR
0	Malaysia	8904	124	8601	179	7	0	1
1	Kenya	17975	285	7833	9857	372	5	90
2	San Marino	699	42	657	0	0	0	0
3	El Salvador	15035	408	7778	6849	405	8	130
4	Gabon	7189	49	4682	2458	205	0	219
...	...	...	...	...	...	...	...	...
2068	Nepal	18752	48	13754	4950	139	3	626
2069	Latvia	1219	31	1045	143	0	0	0
2070	Montenegro	2893	45	809	2039	94	2	70
2071	Bolivia	71181	2647	21478	47056	1752	64	309
2072	Qatar	109597	165	106328	3104	292	0	304

2073 rows × 15 columns

```
datarecon['COUNTRY_WISE_LATEST']['azure-kafka']
```

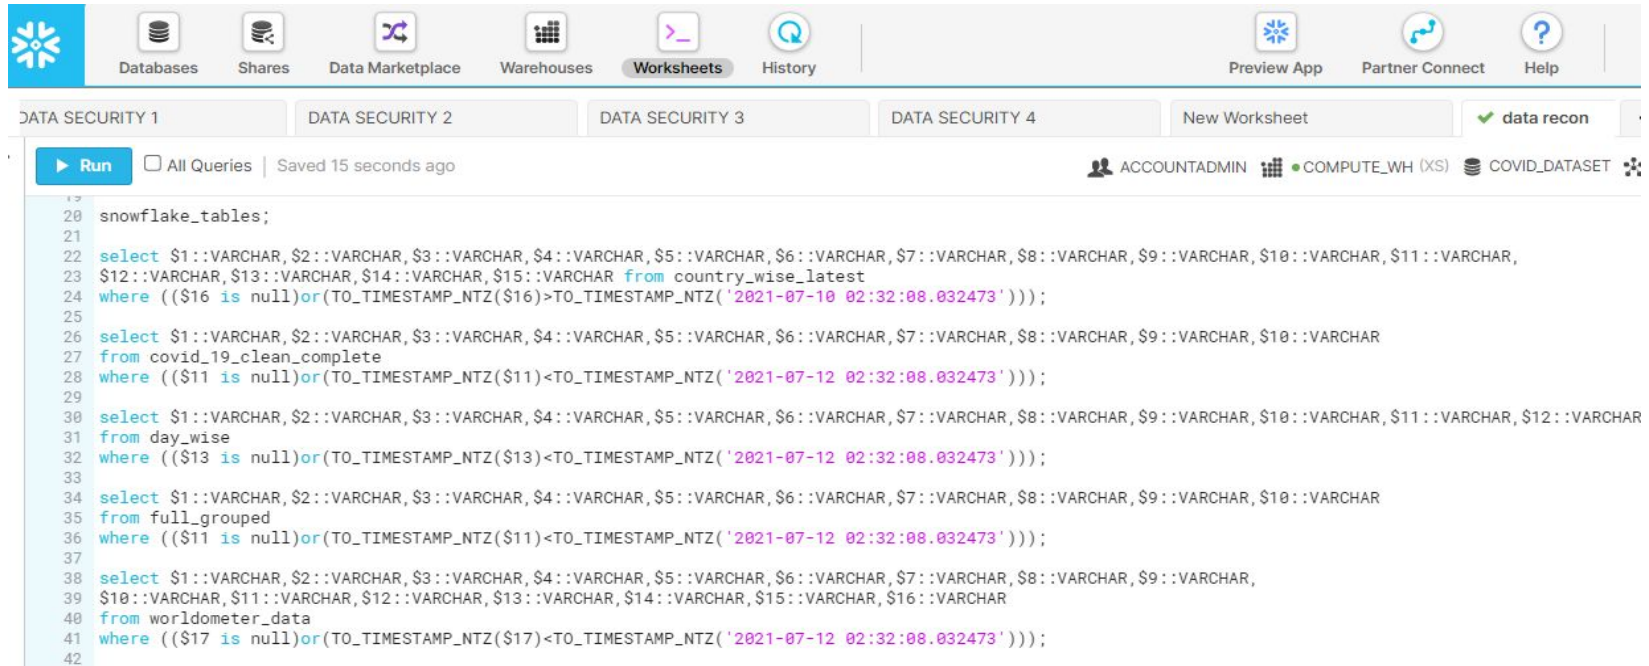
ut[7]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
0	Malaysia	8904	124	8601	179	7	0	1	1.39	96.6	1.44	8800	104	1.18	Western Pacific
1	Malaysia	8904	124	8601	179	7	0	1	1.39	96.6	1.44	8800	104	1.18	Western Pacific
2	Malaysia	8904	124	8601	179	7	0	1	1.39	96.6	1.44	8800	104	1.18	Western Pacific
3	Malaysia	8904	124	8601	179	7	0	1	1.39	96.6	1.44	8800	104	1.18	Western Pacific
4	Malaysia	8904	124	8601	179	7	0	1	1.39	96.6	1.44	8800	104	1.18	Western Pacific
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
1133	Sierra Leone	1783	66	1317	400	0	0	4	3.7	73.86	5.01	1711	72	4.21	Africa
1134	Sierra Leone	1783	66	1317	400	0	0	4	3.7	73.86	5.01	1711	72	4.21	Africa
1135	Sierra Leone	1783	66	1317	400	0	0	4	3.7	73.86	5.01	1711	72	4.21	Africa
1136	Sierra Leone	1783	66	1317	400	0	0	4	3.7	73.86	5.01	1711	72	4.21	Africa
1137	Sierra Leone	1783	66	1317	400	0	0	4	3.7	73.86	5.01	1711	72	4.21	Africa

1138 rows × 15 columns

# Data Reconciliation

Sql code used by python to access table data from snowflake



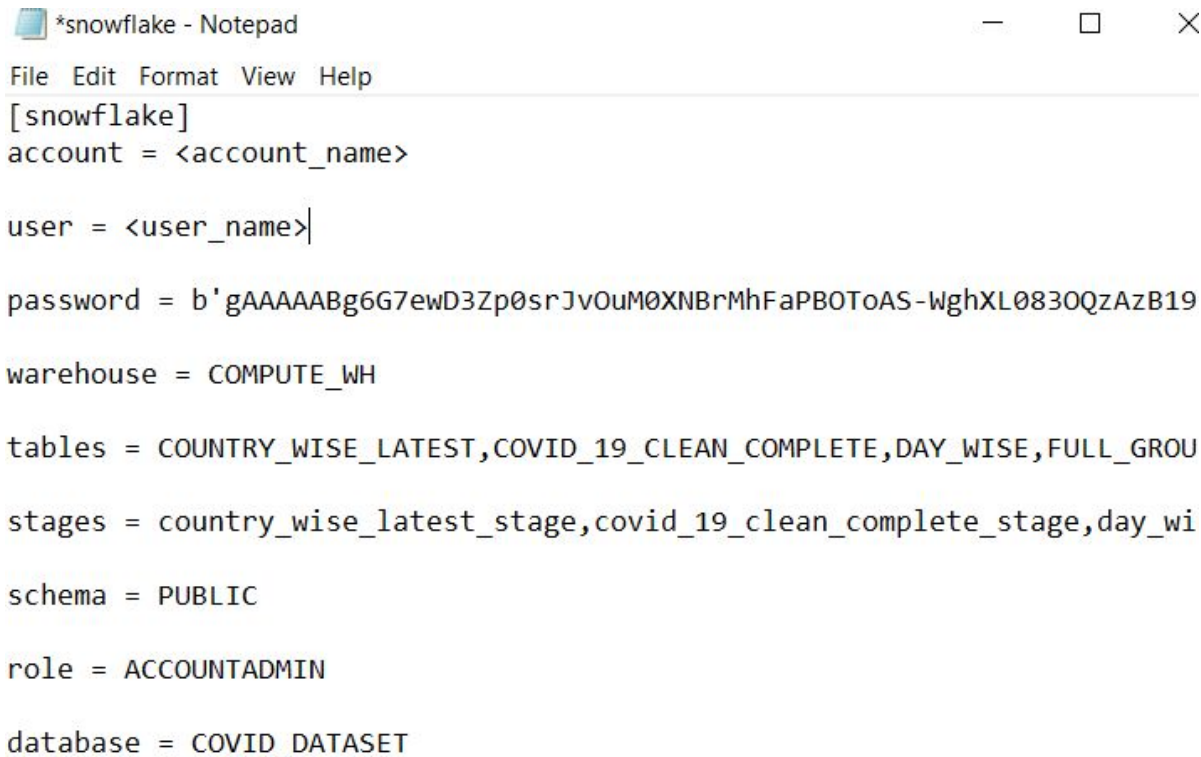
The screenshot displays the Snowflake SQL Editor interface. The top navigation bar includes icons for Databases, Shares, Data Marketplace, Warehouses, Worksheets (active), History, Preview App, Partner Connect, and Help. Below the navigation bar, there are tabs for DATA SECURITY 1, DATA SECURITY 2, DATA SECURITY 3, DATA SECURITY 4, New Worksheet, and a green checkmark icon labeled 'data recon'. The main editor area shows a SQL query with line numbers 19 to 42. The query is a multi-part SELECT statement with several WHERE clauses filtering data based on timestamps and null values. The query is as follows:

```
19 snowflake_tables;
20
21
22 select $1::VARCHAR, $2::VARCHAR, $3::VARCHAR, $4::VARCHAR, $5::VARCHAR, $6::VARCHAR, $7::VARCHAR, $8::VARCHAR, $9::VARCHAR, $10::VARCHAR, $11::VARCHAR,
23 $12::VARCHAR, $13::VARCHAR, $14::VARCHAR, $15::VARCHAR from country_wise_latest
24 where (($16 is null)or(TO_TIMESTAMP_NTZ($16)>TO_TIMESTAMP_NTZ('2021-07-10 02:32:08.032473')));
25
26 select $1::VARCHAR, $2::VARCHAR, $3::VARCHAR, $4::VARCHAR, $5::VARCHAR, $6::VARCHAR, $7::VARCHAR, $8::VARCHAR, $9::VARCHAR, $10::VARCHAR
27 from covid_19_clean_complete
28 where (($11 is null)or(TO_TIMESTAMP_NTZ($11)<TO_TIMESTAMP_NTZ('2021-07-12 02:32:08.032473')));
29
30 select $1::VARCHAR, $2::VARCHAR, $3::VARCHAR, $4::VARCHAR, $5::VARCHAR, $6::VARCHAR, $7::VARCHAR, $8::VARCHAR, $9::VARCHAR, $10::VARCHAR, $11::VARCHAR, $12::VARCHAR
31 from day_wise
32 where (($13 is null)or(TO_TIMESTAMP_NTZ($13)<TO_TIMESTAMP_NTZ('2021-07-12 02:32:08.032473')));
33
34 select $1::VARCHAR, $2::VARCHAR, $3::VARCHAR, $4::VARCHAR, $5::VARCHAR, $6::VARCHAR, $7::VARCHAR, $8::VARCHAR, $9::VARCHAR, $10::VARCHAR
35 from full_grouped
36 where (($11 is null)or(TO_TIMESTAMP_NTZ($11)<TO_TIMESTAMP_NTZ('2021-07-12 02:32:08.032473')));
37
38 select $1::VARCHAR, $2::VARCHAR, $3::VARCHAR, $4::VARCHAR, $5::VARCHAR, $6::VARCHAR, $7::VARCHAR, $8::VARCHAR, $9::VARCHAR,
39 $10::VARCHAR, $11::VARCHAR, $12::VARCHAR, $13::VARCHAR, $14::VARCHAR, $15::VARCHAR, $16::VARCHAR
40 from worldometer_data
41 where (($17 is null)or(TO_TIMESTAMP_NTZ($17)<TO_TIMESTAMP_NTZ('2021-07-12 02:32:08.032473')));
42
```

At the bottom of the editor, there is a 'Run' button and a status bar indicating 'All Queries' and 'Saved 15 seconds ago'. The right side of the interface shows the user 'ACCOUNTADMIN', the warehouse 'COMPUTE\_WH (XS)', and the database 'COVID\_DATASET'.

# Data Reconciliation

Setting snowflake config file to connect python code to snowflake



```
*snowflake - Notepad
File Edit Format View Help
[snowflake]
account = <account_name>

user = <user_name>

password = b'gAAAAABg6G7ewD3Zp0srJvOuM0XNBrMhFaPB0ToAS-WghXL0830QzAzB19

warehouse = COMPUTE_WH

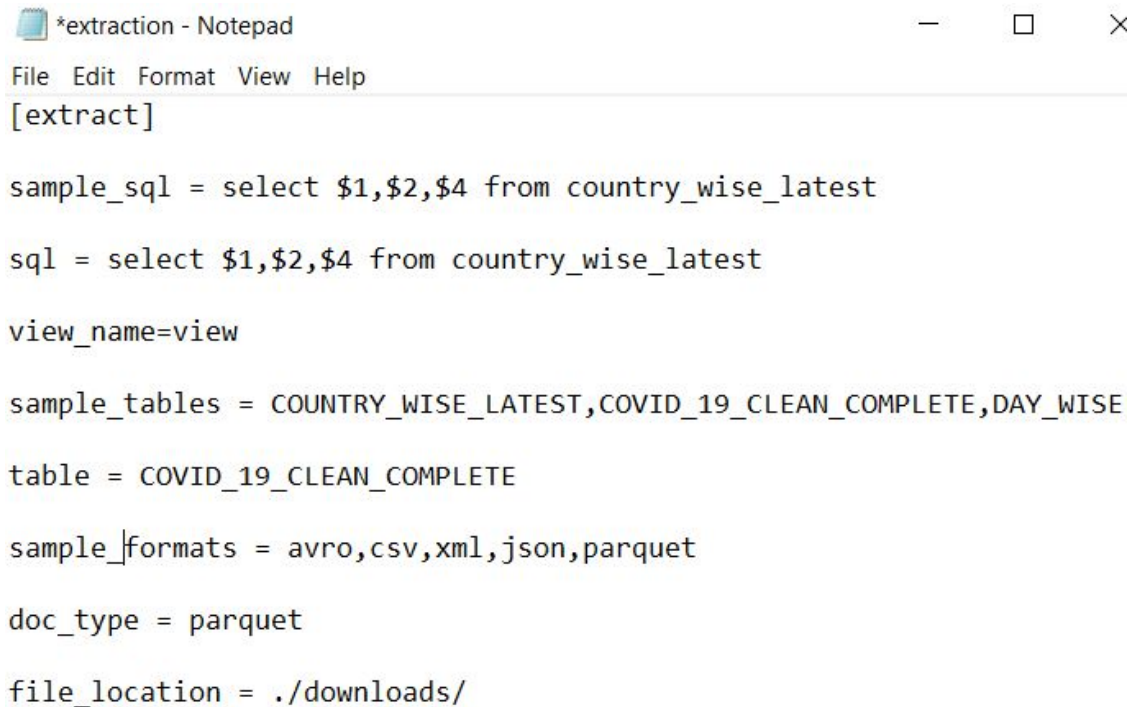
tables = COUNTRY_WISE_LATEST,COVID_19_CLEAN_COMPLETE,DAY_WISE,FULL_GROU
stages = country_wise_latest_stage,covid_19_clean_complete_stage,day_wi
schema = PUBLIC

role = ACCOUNTADMIN

database = COVID_DATASET
```

# Data Extraction

Setting the config file to download a view or table in the required file format



The image shows a Notepad window with the title bar '\*extraction - Notepad'. The menu bar includes 'File', 'Edit', 'Format', 'View', and 'Help'. The text content of the file is as follows:

```
[extract]

sample_sql = select $1,$2,$4 from country_wise_latest

sql = select $1,$2,$4 from country_wise_latest

view_name=view

sample_tables = COUNTRY_WISE_LATEST,COVID_19_CLEAN_COMPLETE,DAY_WISE

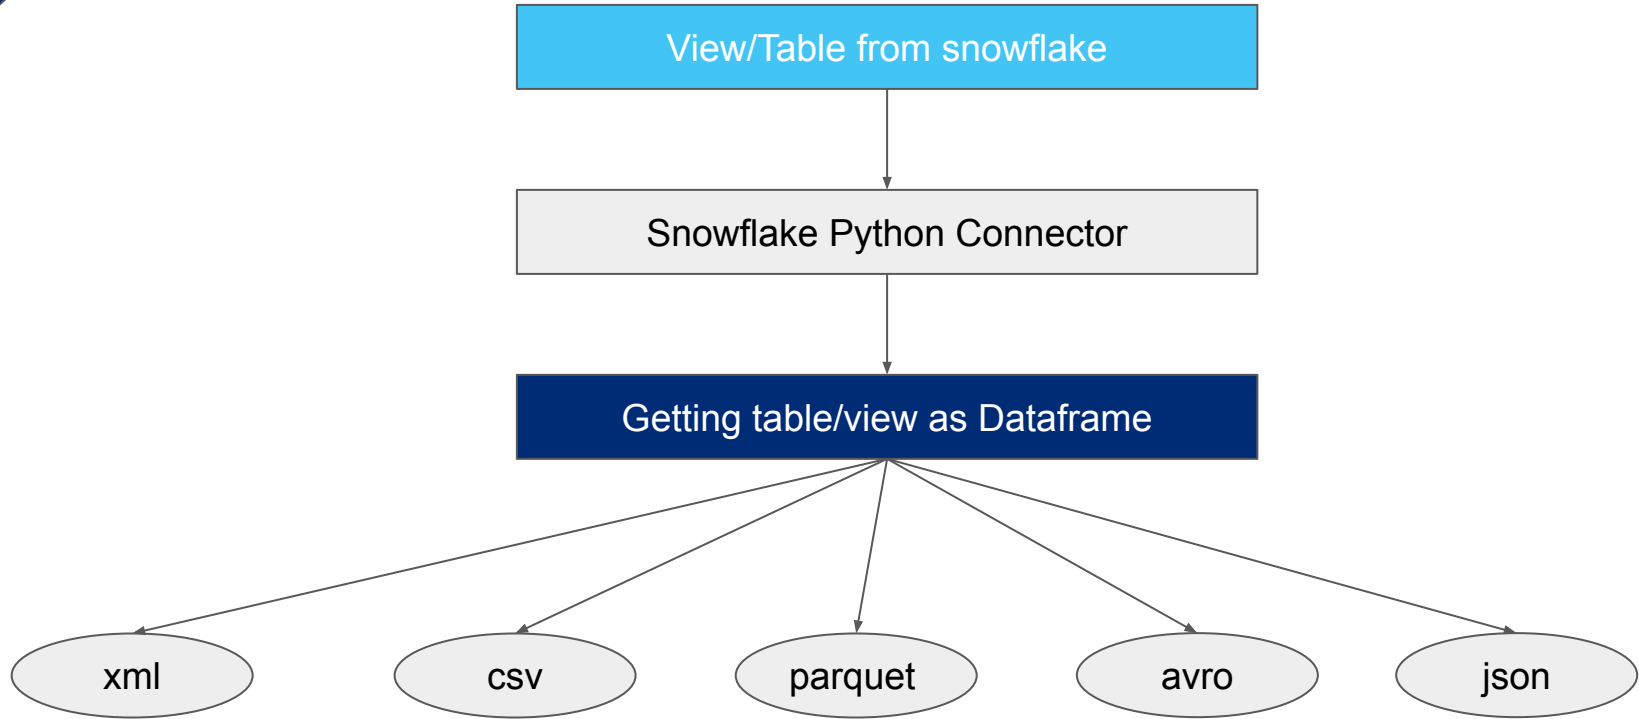
table = COVID_19_CLEAN_COMPLETE

sample_formats = avro, csv, xml, json, parquet

doc_type = parquet

file_location = ./downloads/
```

# Data Extraction



# **Establish Ingestion and Extraction patterns using Azure and Snowflake. (For Batch-time and Real-time)**

**Guided by- Amit Malik**

## **Team Members-**




**Ashwin Nair**



**Rajat Mishra**



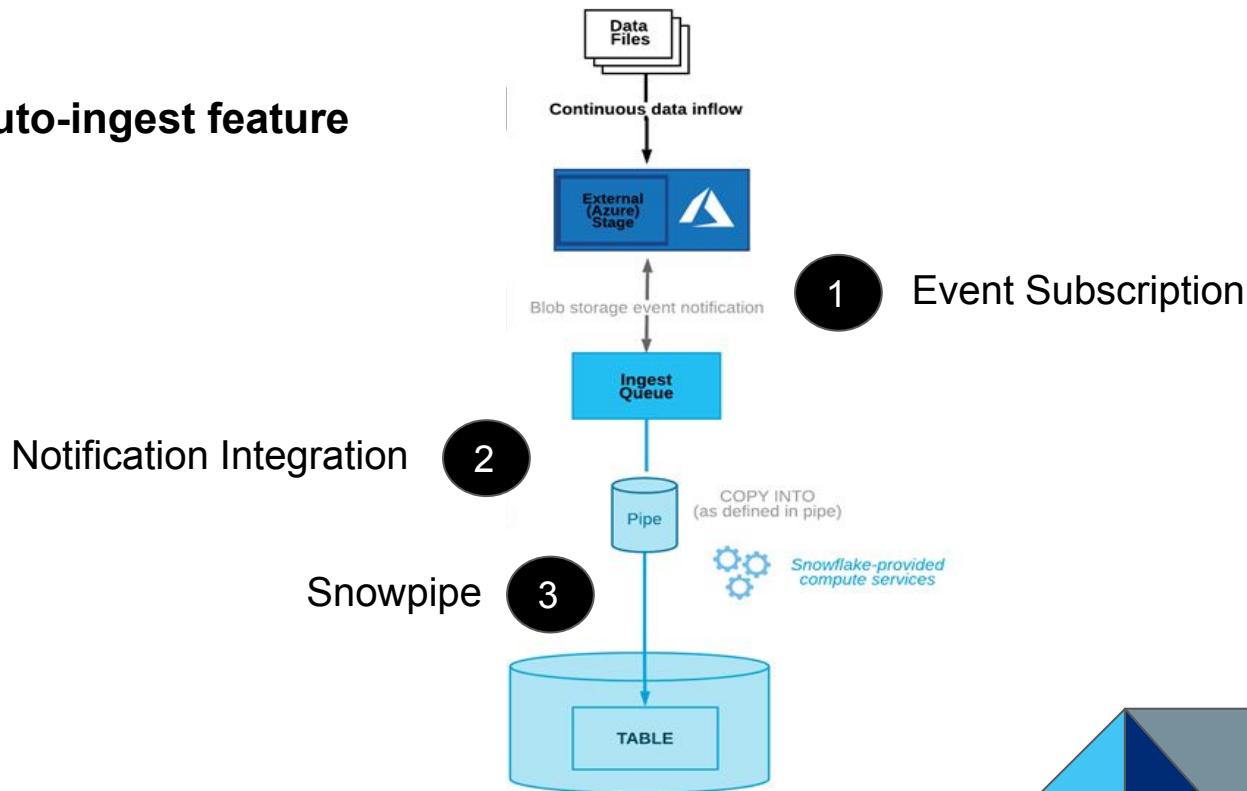
**Nehal Tiwari**






# Data Ingestion

## Snowpipe's auto-ingest feature




# Data Ingestion

- **Connecting Azure Blob Containers with Ingestion Queue via Event Subscription.**




# Data Ingestion

- Connecting the Ingestion Queue with Snowflake Stage(external) via Notification Integration.



# Data Ingestion

- **Snowpipe Automatically loading data into Snowflake Table.**



# Code Snippets


▶ Run ☐ All Queries | Saved 9 seconds ago

ACCOUNTADMIN COMPUTE\_WH (XS) COVID\_DATASET PUBLIC ▾

```
1 //creating notification intergration
2 create or replace notification integration azureintegration
3   enabled = true
4   type = queue
5   notification_provider = azure_storage_queue
6   azure_storage_queue_primary_uri = 'https://ashwinstorage.queue.core.windows.net/ingestionq'
7   azure_tenant_id = '311bd0d6-0163-4fe0-99c5-9c453d609df3';
8
9 //connecting intergration with external stage
10 CREATE OR REPLACE STAGE "COVID_DATASET"."PUBLIC".country_wise_latest_stage URL = 'azure://ashwinstorage.blob.core.windows.net/country-wise-latest-container' CREDENTIALS = (AZURE_SAS_TOKEN = '?sp=racwd1&st=2021-07-08T13:45:00Z&sv=2018-08-01&sr=b&sig=...'
11
12 //Snowpipe's automatic loading data into Snowflake Tables
13 create or replace pipe country_wise_latest_pipe
14   auto_ingest = true
15   integration = azureintegration
16   as
17   copy into "COVID_DATASET"."PUBLIC"."COUNTRY_WISE_LATEST"
18   from
19     (select $1,$2,$3,$4,$5,$6,$7,$8,$9,$10,$11,$12,$13,$14,$15,$16,current_timestamp(3) from @country_wise_latest_stage)
20
21 ;
```

# Data Transformation

- Why Data Transformation?



**Data** is **transformed** to make it better-organized. **Transformed data** would be easier for both humans and computers to use. It acts as a power booster for the **Data analytics process**.

# Data Transformation

- What are Streams and Tasks?

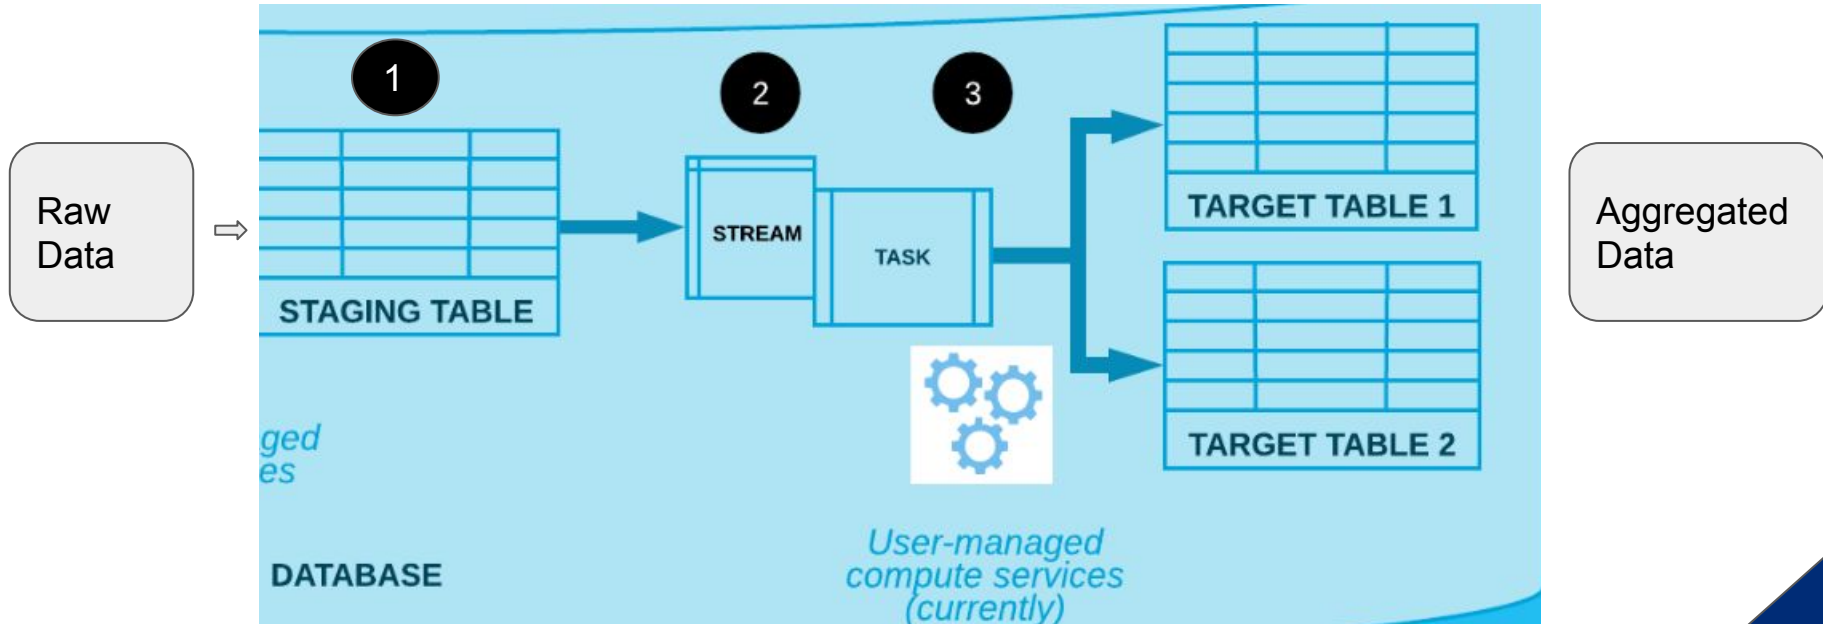


A **stream** is a **Snowflake** object type that provides change data capture (CDC).

A **Task** consumes these streams to run a scheduled query.

# Data Transformation

- How Do Streams and Tasks work?





# Code Snippets

Run

☐ All Queries

Saved 0 seconds ago

ACCOUNTADMIN

COMPUTE\_WH (XS)

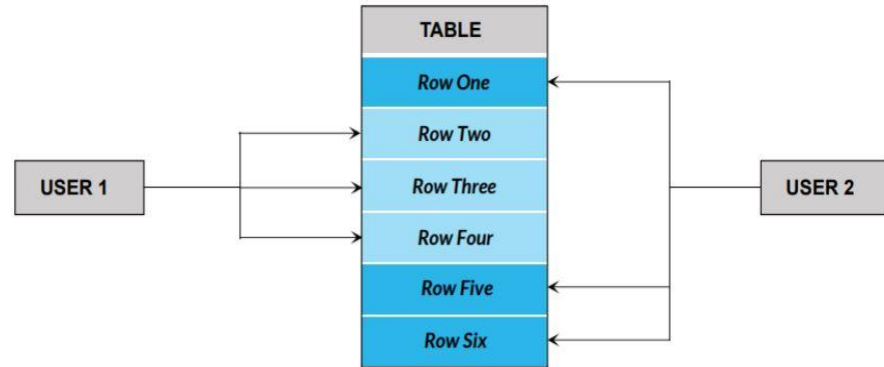
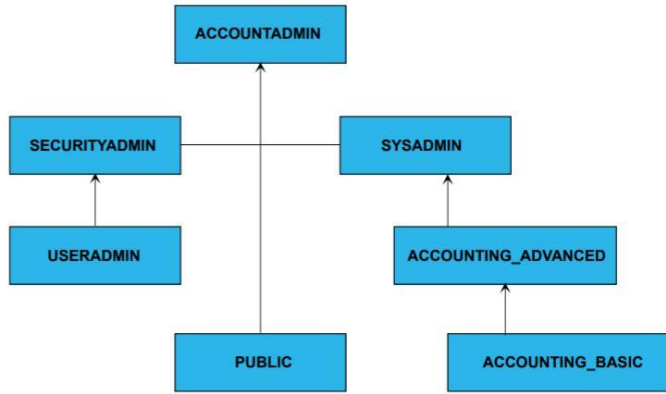
COVID\_DATASET

PUBLIC

..

# Data Security

- Most important aspect of data handling.
- Role level and row level access.
- 4 roles created coherent with the data set used.



**We have created four roles aligning to the dataset  
we have been using:**



**WHO Role**



**PMO India Role**



**BRICS Role**



**SAARC Role**

# Code overview

## 1. Creating Configuration tables:

```
--using role,database,schema
use role accountadmin;
use database covid_dataset;
use schema public;

--create configuration table
create or replace table config_country_table_india( role_name varchar, region varchar);
insert into config_country_table_india values ('PMO_INDIA_CWL','India');

create or replace table config_country_table_brics( role_name varchar, region varchar);
insert into config_country_table_brics values ('BRICS_CWL','India');
insert into config_country_table_brics values ('BRICS_CWL','Brazil');
insert into config_country_table_brics values ('BRICS_CWL','China');
insert into config_country_table_brics values ('BRICS_CWL','Russia');
insert into config_country_table_brics values ('BRICS_CWL','South Africa');

create or replace table config_country_table_saarc( role_name varchar, region varchar);
insert into config_country_table_saarc values ('SAARC_CWL','India');
insert into config_country_table_saarc values ('SAARC_CWL','Nepal');
insert into config_country_table_saarc values ('SAARC_CWL','Pakistan');
insert into config_country_table_saarc values ('SAARC_CWL','Bangladesh');
insert into config_country_table_saarc values ('SAARC_CWL','Afghanistan');
insert into config_country_table_saarc values ('SAARC_CWL','Maldives');
insert into config_country_table_saarc values ('SAARC_CWL','Bhutan');
insert into config_country_table_saarc values ('SAARC_CWL','Sri Lanka');
```

## 2. Creating roles:

```
--create roles
create role pmo_india_cwl;
create role brics_cwl;
create role saarc_cwl;
create role who_cwl;
```

### 3. Creating secure views:

```
--create secure view
create or replace secure view sv_country_wise_latest_india as
select *
from country_wise_latest
where "Country/Region" in (
    select region from config_country_table_india
    where role_name = CURRENT_ROLE()
);

create or replace secure view sv_country_wise_latest_brics as
select *
from country_wise_latest
where "Country/Region" in (
    select region from config_country_table_brics
    where role_name = CURRENT_ROLE()
);

create or replace secure view sv_country_wise_latest_saarc as
select *
from country_wise_latest
where "Country/Region" in (
    select region from config_country_table_saarc
    where role_name = CURRENT_ROLE()
);

create or replace secure view sv_country_wise_latest_who as
select *
from country_wise_latest
;
```

### 4. Granting permissions for usage on warehouse, database, schema and secure view:

```
--grant permissions"COVID_DATASET"."PUBLIC"."COUNTRY_WISE_LATEST"
use role sysadmin;
grant usage on warehouse compute_wh to role pmo_india_cwl;
grant usage on warehouse compute_wh to role brics_cwl;
grant usage on warehouse compute_wh to role saarc_cwl;
grant usage on warehouse compute_wh to role who_cwl;

grant usage on database covid_dataset to role pmo_india_cwl;
grant usage on database covid_dataset to role brics_cwl;
grant usage on database covid_dataset to role saarc_cwl;
grant usage on database covid_dataset to role who_cwl;

grant usage on schema public to role pmo_india_cwl;
grant usage on schema public to role brics_cwl;
grant usage on schema public to role saarc_cwl;
grant usage on schema public to role who_cwl;

--grant view
use role accountadmin;
grant select on view sv_country_wise_latest_india to role pmo_india_cwl ;
grant select on view sv_country_wise_latest_brics to role brics_cwl ;
grant select on view sv_country_wise_latest_saarc to role saarc_cwl ;
grant select on view sv_country_wise_latest_who to role who_cwl ;
```

## 5. Creating usernames and passwords for each roles:

```
--create username and password
create or replace user pmo_india_cwl PASSWORD = 'ind001' COMMENT='created this user to check row level security'
LOGIN_NAME='ind_cwl_user' DISPLAY_NAME='INDCWL' DEFAULT_ROLE='PUBLIC'
MUST_CHANGE_PASSWORD= FALSE;

create or replace user brics_cwl PASSWORD = 'brics001' COMMENT='created this user to check row level security'
LOGIN_NAME='brics_cwl_user' DISPLAY_NAME='BRICSCWL' DEFAULT_ROLE='PUBLIC'
MUST_CHANGE_PASSWORD= FALSE;

create or replace user saarc_cwl PASSWORD = 'saarc001' COMMENT='created this user to check row level security'
LOGIN_NAME='saarc_cwl_user' DISPLAY_NAME='SAARCCWL' DEFAULT_ROLE='PUBLIC'
MUST_CHANGE_PASSWORD= FALSE;

create or replace user who_cwl PASSWORD = 'who001' COMMENT='created this user to check row level security'
LOGIN_NAME='who_cwl_user' DISPLAY_NAME='WHOCWL' DEFAULT_ROLE='PUBLIC'
MUST_CHANGE_PASSWORD= FALSE;
```

## 6. Granting roles to users for login in secure views:

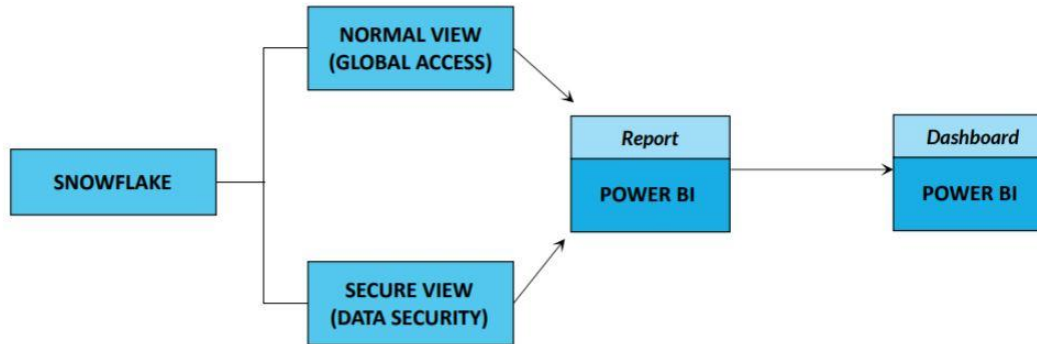
```
--grant role
grant role pmo_india_cwl to user pmo_india_cwl ;
grant role brics_cwl to user brics_cwl ;
grant role saarc_cwl to user saarc_cwl ;
grant role who_cwl to user who_cwl;
-----
--new worksheet
select current_role();
use role pmo_india_cwl;
select * from covid_dataset.public.sv_country_wise_latest_india;

select current_role();
use role brics_cwl;
select * from covid_dataset.public.sv_country_wise_latest_brics;

select current_role();
use role saarc_cwl;
select * from covid_dataset.public.sv_country_wise_latest_saarc;

select current_role();
use role who_cwl;
select * from covid_dataset.public.sv_country_wise_latest_who;
```

# Data Visualisation



- Dashboard for visual representation of data using Power BI.
- Facilitated comparison between countries.
- Visualisation helps in understanding the dataset better.



# Takeaways

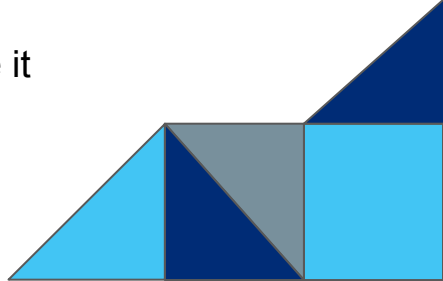
## **Rajat Mishra :**

- Never Hesitate to ask your doubts with your seniors/mentors/peers. Clarify it before its late.
- Learning your team members and synchronizing with them is a key aspect.

## **Ashwin Nair :**

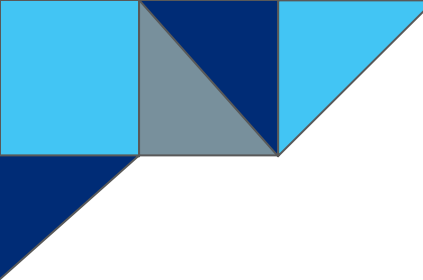
- Identifying core strengths of team members helps in overall project execution
- Helping your team mates boosts team coordination and team efficiency

## **Nehal Tiwari :**

- No matter how unexplored field is, there is always a way to ace it by learning.
  - Coordinating between members of team for a better finished product.
- 







**THANK YOU!!**

